



先进编译实验室
Advanced Compiler

先进编译实验室
Advanced Compiler

Triton程序编写

先进编译实验室
Advanced Compiler

嘉宾：杜玉博



先进编译实验室
Advanced Compiler

常见的Triton API语法

triton

<code>jit</code>	Decorator for JIT-compiling a function using the Triton compiler.
<code>autotune</code>	Decorator for auto-tuning a <code>triton.jit</code> 'd function.
<code>heuristics</code>	Decorator for specifying how the values of certain meta-parameters may be computed.
<code>config</code>	An object that represents a possible kernel configuration for the auto-tuner to try.



triton.autotune

```
triton.autotune(configs, key, prune_configs_by=None, reset_to_zero=None, restore_value=None,  
pre_hook=None, post_hook=None, warmup=25, rep=100, use_cuda_graph=False)
```

Decorator for auto-tuning a `triton.jit`'d function.

```
@triton.autotune(configs=[  
    triton.Config(kwarg={'BLOCK_SIZE': 128}, num_warps=4),  
    triton.Config(kwarg={'BLOCK_SIZE': 1024}, num_warps=8),  
],  
key=['x_size'] # the two above configs will be evaluated anytime  
               # the value of x_size changes  
)  
@triton.jit  
def kernel(x_ptr, x_size, **META):  
    BLOCK_SIZE = META['BLOCK_SIZE']
```

Note: When all the configurations are evaluated, the kernel will run multiple times. This means that whatever value the kernel updates will be updated multiple times. To avoid this undesired behavior, you can use the `reset_to_zero` argument, which resets the value of the provided tensor to zero before running any configuration.





triton.Config

```
class triton.Config(self, kwargs, num_warps=4, num_stages=2, num_ctas=1, maxnreg=None,  
pre_hook=None)
```

An object that represents a possible kernel configuration for the auto-tuner to try.

- Variables:
- **kwargs** – a dictionary of meta-parameters to pass to the kernel as keyword arguments.
 - **num_warps** – the number of warps to use for the kernel when compiled for GPUs. For example, if `num_warps=8`, then each kernel instance will be automatically parallelized to cooperatively execute using $8 * 32 = 256$ threads.
 - **num_stages** – the number of stages that the compiler should use when software-pipelining loops. Mostly useful for matrix multiplication workloads on SM80+ GPUs.
 - **num_ctas** – number of blocks in a block cluster. SM90+ only.
 - **maxnreg** – maximum number of registers one thread can use. Corresponds to `ptx .maxnreg` directive. Not supported on all platforms.
 - **pre_hook** – a function that will be called before the kernel is called. Parameters of this function are args.



常见的Triton API语法

Math Ops

<code>abs</code>	Computes the element-wise absolute value of <code>x</code> .
<code>cdiv</code>	Computes the ceiling division of <code>x</code> by <code>div</code> .
<code>ceil</code>	Computes the element-wise ceil of <code>x</code> .
<code>cos</code>	Computes the element-wise cosine of <code>x</code> .
<code>sin</code>	Computes the element-wise sine of <code>x</code> .
<code>softmax</code>	Computes the element-wise softmax of <code>x</code> .
<code>sqrt</code>	Computes the element-wise fast square root of <code>x</code> .



常见的Triton API语法

Debug Ops

<code>static_print</code>	Print the values at compile time.
<code>static_assert</code>	Assert the condition at compile time.
<code>device_print</code>	Print the values at runtime from the device.
<code>device_assert</code>	Assert the condition at runtime from the device.





向量加法示例

```
@triton.jit
def add_kernel(x_ptr, y_ptr, output_ptr, n_elements, BLOCK_SIZE: tl.constexpr, ):
    pid = tl.program_id(axis=0)
    block_start = pid * BLOCK_SIZE
    offsets = block_start + tl.arange(0, BLOCK_SIZE)
    mask = offsets < n_elements
    x = tl.load(x_ptr + offsets, mask=mask)
    y = tl.load(y_ptr + offsets, mask=mask)
    output = x + y
    tl.store(output_ptr + offsets, output, mask=mask)
```





如何调用内核函数

```
def add(x: torch.Tensor, y: torch.Tensor):  
    output = torch.empty_like(x)  
    assert x.is_cuda and y.is_cuda and output.is_cuda  
    n_elements = output.numel()  
    grid = lambda meta: (triton.cdiv(n_elements, meta['BLOCK_SIZE']), )  
    add_kernel[grid](x, y, output, n_elements, BLOCK_SIZE=1024)  
    return output
```





编译运行过程：

```
# python 01-vector-add.py
```

运行结果分析：

```
(triton2.1.0) [duyb@Stream9 tutorials]$ python 01-vector-add.py
tensor([1.3713, 1.3076, 0.4940, ..., 1.3374, 1.4960, 0.9115], device='cuda:0')
tensor([1.3713, 1.3076, 0.4940, ..., 1.3374, 1.4960, 0.9115], device='cuda:0')
The maximum difference between torch and triton is 0.0
/public/home/duyb/anaconda3/envs/triton2.1.0/lib/python3.11/site-packages/triton/testing.py:304: UserWarning: FigureCanvasAgg is non-interactive, and thus cannot be shown
  plt.show()
vector-add-performance:

```

	size	Triton	Torch
0	4096.0	15.999999	15.999999
1	8192.0	24.000000	24.000000
2	16384.0	48.000000	48.000000
3	32768.0	90.352949	94.523077
4	65536.0	127.999995	127.999995
5	131072.0	192.000000	192.000000
6	262144.0	255.999991	250.137408
7	524288.0	279.272725	279.272725
8	1048576.0	307.200008	307.200008
9	2097152.0	319.168844	322.440352
10	4194304.0	327.714130	329.775424
11	8388608.0	332.108094	332.353736
12	16777216.0	334.367350	334.376254
13	33554432.0	335.508544	335.526441
14	67108864.0	335.938503	336.082067
15	134217728.0	336.225721	336.369531



AdvancedCompiler

Tel: 13839830713