



Triton-IR剖析 (上)

嘉宾：董贞汝

Triton kernel到ttir转换过程



以01-vector_add.py为例简单了解转换过程

调用Triton Kernel (python/tutorials/01-vector-add.py) : `add(x, y)`



JIT调用编译器(triton/runtime/jit.py): `kernel = self.compile()`



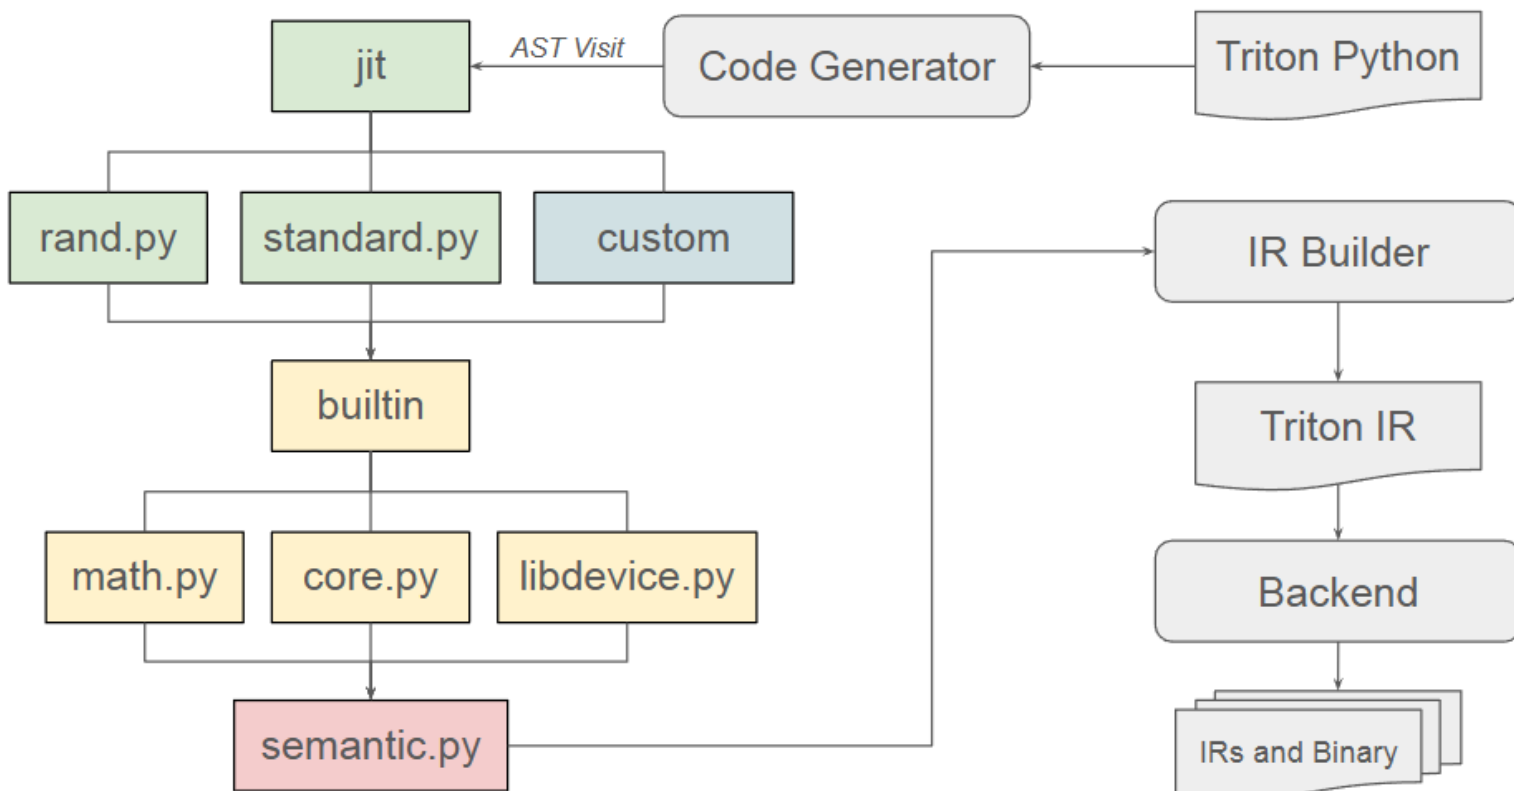
调用make_ir(triton/compiler/compiler.py): `src.make_ir()`



调用ast去生成ttir(triton/compiler/compiler.py): `ast_to_ttir()`, 逻辑位于triton/compiler/code_generator.py文件中,
具体地进行codegen的遍历完成Triton Python到Triton IR (ttir) 的转换



Triton kernel到ttir转换过程



- triton/language/random.py
随机数生成
- triton/language/standard.py
一些算子的标准实现
- triton/language/math.py
数学算子的实现
- triton/language/core.py
Triton核心操作的实现
- triton/language/extra/cuda/libdevice.py
硬件相关的操作实现
- triton/language/semantic.py
语法



kernel与Triton IR的对应



```
1 import triton
2 import triton.language as tl
3
4 @triton.jit
5 def add_kernel(
6     x_ptr, # *Pointer* to first input vector.
7     y_ptr, # *Pointer* to second input vector.
8     output_ptr, # *Pointer* to output vector.
9     n_elements, # Size of the vector.
10     BLOCK_SIZE: tl.constexpr, # Number of elements each program should process.
11     # NOTE: 'constexpr' so it can be used as a shape value.
12 ):
13     # There are multiple 'programs' processing different data. We identify which program
14     # we are here:
15     pid = tl.program_id(axis=0) # We use a 1D launch grid so axis is 0.
16     # This program will process inputs that are offset from the initial data.
17     # For instance, if you had a vector of length 256 and block_size of 64, the programs
18     # would each access the elements [0:64, 64:128, 128:192, 192:256].
19     # Note that offsets is a list of pointers:
20     block_start = pid * BLOCK_SIZE
21     offsets = block_start + tl.arange(0, BLOCK_SIZE)
22     # Create a mask to guard memory operations against out-of-bounds accesses.
23     mask = offsets < n_elements
24     # Load x and y from DRAM, masking out any extra elements in case the input is not a
25     # multiple of the block size.
26     x = tl.load(x_ptr + offsets, mask=mask)
27     y = tl.load(y_ptr + offsets, mask=mask)
28     output = x + y
29     # Write x + y back to DRAM.
30     tl.store(output_ptr + offsets, output, mask=mask)
31
```



```
3 module {
4     tt.func public @add_kernel_01234(%arg0: !tt.ptr<f32>, %arg1: !tt.ptr<f32>, %arg2: !tt.ptr<f32>, %arg3: i32) {
5         %c1024_i32 = arith.constant 1024 : i32
6         %0 = tt.get_program_id {axis = 0 : i32} : i32
7         %1 = arith.muli %0, %c1024_i32 : i32
8         %2 = tt.make_range {end = 1024 : i32, start = 0 : i32} : tensor<1024xi32>
9         %3 = tt.splat %1 : (i32) -> tensor<1024xi32>
10        %4 = arith.addi %3, %2 : tensor<1024xi32>
11        %5 = tt.splat %arg3 : (i32) -> tensor<1024xi32>
12        %6 = arith.cmpi slt, %4, %5 : tensor<1024xi32>
13        %7 = tt.splat %arg0 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
14        %8 = tt.addptr %7, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
15        %9 = tt.load %8, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile = false} : tensor<1024xf32>
16        %10 = tt.splat %arg1 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
17        %11 = tt.addptr %10, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
18        %12 = tt.load %11, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile = false} : tensor<1024xf32>
19        %13 = arith.addf %9, %12 : tensor<1024xf32>
20        %14 = tt.splat %arg2 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
21        %15 = tt.addptr %14, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
22        tt.store %15, %13, %6 {cache = 1 : i32, evict = 1 : i32} : tensor<1024xf32>
23        tt.return
24    }
25 }
```



IR Structure at a Glance

Result name

Op name

Operands and result type

```
%2 = arith.addf %0, %1 : tensor<1024xf32>
```

Dialect prefix

Operands

- MLIR uses the notion of operations
- Operations are “opaque functions” to MLIR
- IR is represented in SSA form
- Refer to these [slides](#) for more comprehensive examples

IR Structure at a Glance

- Dialects
 - Collection of operations, types, attributes, related transformation passes, analysis, etc.
 - Multiple dialects can co-exist within one module
 - E.g.: arith, math, func, llvm, memref, etc.
- Types
 - Supports both builtin and custom Dialect types
 - E.g.: i8, f32, memref, vector, tensor, etc.
- Attributes
 - In-place constant data on operations
 - E.g., comparison predicate, if perform boundary check on triton load, etc.



先进编译实验室
Advanced Compiler


Triton-IR剖析 (下)

嘉宾：董贞汝



先进编译实验室
Advanced Compiler

Triton IR中的优化pass



TritonTypes.td:

- Float types
- Integer types
- Tensor types
- Pointer types
- ...

TritonOps.td:

- tt.addptr
- tt.load
- tt.store
- tt.dot
- ...

Passes.td:

- ...

TritonTypeInterfaces.td

TritonInterfaces.td

TritonAttrDefs.td

- Eviction policies
- Padding options
- Atomics attrs
- ...

编译过程

```
triton-main
├── include
├── triton
│   ├── Conversion
│   ├── Dialect
│   │   ├── Triton
│   │   └── IR
│       ├── CMakeLists.txt
│       ├── Dialect.h
│       ├── Interfaces.h
│       ├── Traits.h
│       ├── TritonAttrDefs.td
│       ├── TritonDialect.td
│       ├── TritonInterfaces.td
│       ├── TritonOps.td
│       ├── TritonTypeInterfaces.td
│       ├── TritonTypes.td
│       ├── Types.h
│       ├── Utility.h
│       ├── Transforms
│       ├── CMakeLists.txt
│       ├── Passes.h
│       └── Passes.td
│       ├── CMakeLists.txt
│       ├── TritonGPU
│       ├── TritonNvidiaGPU
│       └── CMakeLists.txt
```



Triton IR中的优化pass



在 `python/triton/compiler/compiler.py` 调用 `make_ttir` 去运行一些优化Pass。

```
for ext, compile_ir in list(stages.items())[first_stage:]:  
    next_module = compile_ir(module, metadata)
```

NVIDIA GPU对应的`make_ttir`优化pass pipeline位于`triton/backends/nvidia/compiler.py`中

```
@staticmethod  
def make_ttir(mod, metadata, opt):  
    breakpoint()  
    pm = ir.pass_manager(mod.context)  
    pm.enable_debug()  
    passes.common.add_inliner(pm)  
    passes.ttir.add_rewrite_tensor_pointer(pm)  
    passes.ttir.add_combine(pm)  
    passes.common.add_canonicalizer(pm)  
    passes.ttir.add_reorder_broadcast(pm)  
    passes.common.add_cse(pm)  
    passes.common.add_lcm(pm)  
    passes.common.add_symbol_dce(pm)  
    pm.run(mod)  
    return mod
```

Triton IR中的优化pass



Cse Pass: 公共子表达式消除 (CSE, Common Subexpression Elimination)

功能: 识别和消除重复计算的表达式, 优化程序的运行时间和空间。

执行命令: `triton-opt cse_before.ttir -cse &> cse_after.ttir`

```
1 module {
2   tt.func public @add_kernel_0d1d2d3d(%arg0: !tt.ptr<f32> {tt.divisibility = 16 : i32},
3     %c1024_i32 = arith.constant 1024 : i32
4     %0 = tt.get_program_id x : i32
5     %1 = arith.muli %0, %c1024_i32 : i32
6     %2 = tt.make_range {end = 1024 : i32, start = 0 : i32} : tensor<1024xi32>
7     %3 = tt.splat %1 : (i32) -> tensor<1024xi32>
8     %4 = arith.addi %3, %2 : tensor<1024xi32>
9     %5 = tt.splat %arg3 : (i32) -> tensor<1024xi32>
10    %6 = arith.cmpi slt, %4, %5 : tensor<1024xi32>
11    %7 = tt.splat %arg0 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
12    %8 = tt.addptr %7, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
13    %9 = tt.load %8, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile = false} : tensor<1024xf32>
14    %10 = tt.splat %arg1 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
15    %11 = tt.addptr %10, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
16    %12 = tt.load %11, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile = false} : tensor<1024xf32>
17    %13 = arith.addf %9, %12 : tensor<1024xf32>
18    %14 = tt.splat %arg2 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
19    %15 = tt.addptr %14, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
20    %16 = tt.addptr %14, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
21    tt.store %15, %13, %6 {cache = 1 : i32, evict = 1 : i32} : tensor<1024xf32>
22    tt.return
23  }
24 }
```

```
1 module {
2   tt.func public @add_kernel_0d1d2d3d(%arg0: !tt.ptr<f32> {tt.divisibility = 16 : i32},
3     %c1024_i32 = arith.constant 1024 : i32
4     %0 = tt.get_program_id x : i32
5     %1 = arith.muli %0, %c1024_i32 : i32
6     %2 = tt.make_range {end = 1024 : i32, start = 0 : i32} : tensor<1024xi32>
7     %3 = tt.splat %1 : (i32) -> tensor<1024xi32>
8     %4 = arith.addi %3, %2 : tensor<1024xi32>
9     %5 = tt.splat %arg3 : (i32) -> tensor<1024xi32>
10    %6 = arith.cmpi slt, %4, %5 : tensor<1024xi32>
11    %7 = tt.splat %arg0 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
12    %8 = tt.addptr %7, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
13    %9 = tt.load %8, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile = false} : tensor<1024xf32>
14    %10 = tt.splat %arg1 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
15    %11 = tt.addptr %10, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
16    %12 = tt.load %11, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile = false} : tensor<1024xf32>
17    %13 = arith.addf %9, %12 : tensor<1024xf32>
18    %14 = tt.splat %arg2 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
19    %15 = tt.addptr %14, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
20    tt.store %15, %13, %6 {cache = 1 : i32, evict = 1 : i32} : tensor<1024xf32>
21    tt.return
22  }
23 }
```

Triton IR中的优化pass



Canonicalizer Pass: 规范化Pass

功能：将代码转换为一种规范化的形式，消除冗余和不必要的复杂结构，以便为后续的优化 Pass 提供便利。

执行命令：triton-opt canonicalizer_before.ttir -canonicalize &> canonicalizer_after.ttir

```
1 module {
2   tt.func public @add_kernel_0d1d2d3d(%arg0: !tt.ptr<f32> {tt.divisibility = 16 : i32}, %a
3     %c1024_i32 = arith.constant 1024 : i32
4     %cst_0 = arith.constant 0 : i32
5     %cst_1 = arith.constant 1 : i32
6     %16 = arith.muli %c1024_i32, %cst_1 : i32
7     %17 = arith.addi %c1024_i32, %cst_0 : i32
8     %0 = tt.get_program_id x : i32
9     %1 = arith.muli %0, %c1024_i32 : i32
10    %2 = tt.make_range {end = 1024 : i32, start = 0 : i32} : tensor<1024xi32>
11    %3 = tt.splat %1 : (i32) -> tensor<1024xi32>
12    %4 = arith.addi %3, %2 : tensor<1024xi32>
13    %5 = tt.splat %arg3 : (i32) -> tensor<1024xi32>
14    %6 = arith.cmpi slt, %4, %5 : tensor<1024xi32>
15    %7 = tt.splat %arg0 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
16    %8 = tt.addptr %7, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
17    %9 = tt.load %8, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile = false} : tensor<10
18    %10 = tt.splat %arg1 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
19    %11 = tt.addptr %10, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
20    %12 = tt.load %11, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile = false} : tensor<
21    %13 = arith.addf %9, %12 : tensor<1024xf32>
22    %14 = tt.splat %arg2 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
23    %15 = tt.addptr %14, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
24    tt.store %15, %13, %6 {cache = 1 : i32, evict = 1 : i32} : tensor<1024xf32>
25    tt.return
26  }
27 }
```

```
1 module {
2   tt.func public @add_kernel_0d1d2d3d(%arg0: !tt.ptr<f32> {tt.divisibility = 16 : i32}, %arg1:
3     %c1024_i32 = arith.constant 1024 : i32
4     %0 = tt.get_program_id x : i32
5     %1 = arith.muli %0, %c1024_i32 : i32
6     %2 = tt.make_range {end = 1024 : i32, start = 0 : i32} : tensor<1024xi32>
7     %3 = tt.splat %1 : (i32) -> tensor<1024xi32>
8     %4 = arith.addi %3, %2 : tensor<1024xi32>
9     %5 = tt.splat %arg3 : (i32) -> tensor<1024xi32>
10    %6 = arith.cmpi slt, %4, %5 : tensor<1024xi32>
11    %7 = tt.splat %arg0 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
12    %8 = tt.addptr %7, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
13    %9 = tt.load %8, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile = false} : tensor<1024xf
14    %10 = tt.splat %arg1 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
15    %11 = tt.addptr %10, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
16    %12 = tt.load %11, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile = false} : tensor<1024
17    %13 = arith.addf %9, %12 : tensor<1024xf32>
18    %14 = tt.splat %arg2 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
19    %15 = tt.addptr %14, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
20    tt.store %15, %13, %6 {cache = 1 : i32, evict = 1 : i32} : tensor<1024xf32>
21    tt.return
22  }
23 }
```

Triton IR中的优化pass



先进编译实验室
Advanced Compiler

12

执行命令: `triton-opt inliner_before.ttir -inline &> inliner_after.ttir`

Inliner Pass:

内联Pass

功能: 直接将函数调用替换为函数体的代码, 避免了函数调用所带来的开销, 从而提高了程序的执行效率。



先进编译实验

Advanced Compiler

```
1 module {
2-   tt.func public @add_kernel_0d1d2d3(%arg0: !tt.ptr<f32> {tt.divisibil
3       %0 = tt.get_program_id x : i32
4-   %c1024_i32 = arith.constant 1024 : i32
5       %1 = arith.muli %0, %c1024_i32 : i32
6       %2 = tt.make_range {end = 1024 : i32, start = 0 : i32} : tensor<10
7       %3 = tt.splat %1 : (i32) -> tensor<1024xi32>
8       %4 = arith.addi %3, %2 : tensor<1024xi32>
9       %5 = tt.splat %arg3 : (i32) -> tensor<1024xi32>
10      %6 = arith.cmpi slt, %4, %5 : tensor<1024xi32>
11-   %7 = tt.make_range {end = 1024 : i32, start = 0 : i32} : tensor<10
12-   %8 = tt.splat %1 : (i32) -> tensor<1024xi32>
13-   %9 = arith.addi %8, %7 : tensor<1024xi32>
14-   %10 = tt.splat %arg3 : (i32) -> tensor<1024xi32>
15-   %11 = arith.cmpi slt, %9, %10 : tensor<1024xi32>
16-   %12 = tt.splat %arg0 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
17-   %13 = tt.addptr %12, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi
18-   %14 = tt.load %13, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile
19-   %15 = tt.splat %arg1 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
20-   %16 = tt.addptr %15, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi
21-   %17 = tt.load %16, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile
22-   %18 = tt.call @add_operator_fp32S1024S_fp32S1024S__(%14, %17) : (t
23-   %19 = tt.splat %arg2 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
24-   %20 = tt.addptr %19, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi
25-   tt.store %20, %18, %6 {cache = 1 : i32, evict = 1 : i32} : tensor<
26-   tt.return
27- }
28- tt.func private @add_operator__fp32S1024S_fp32S1024S__(%arg0: tensor<
29-   %0 = arith.addf %arg0, %arg1 : tensor<1024xf32>
30-   tt.return %0 : tensor<1024xf32>
31- }
```

```
1 module {
2+   tt.func public @add_kernel_0d1d2d3(%arg0: !tt.ptr<f32> {tt.divisibility =
3+   %c1024_i32 = arith.constant 1024 : i32
4       %0 = tt.get_program_id x : i32
5       %1 = arith.muli %0, %c1024_i32 : i32
6       %2 = tt.make_range {end = 1024 : i32, start = 0 : i32} : tensor<1024xi3
7       %3 = tt.splat %1 : (i32) -> tensor<1024xi32>
8       %4 = arith.addi %3, %2 : tensor<1024xi32>
9       %5 = tt.splat %arg3 : (i32) -> tensor<1024xi32>
10      %6 = arith.cmpi slt, %4, %5 : tensor<1024xi32>
11+   %7 = tt.splat %arg0 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
12+   %8 = tt.addptr %7, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
13+   %9 = tt.load %8, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile = fal
14+   %10 = tt.splat %arg1 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
15+   %11 = tt.addptr %10, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
16+   %12 = tt.load %11, %6 {cache = 1 : i32, evict = 1 : i32, isVolatile = f
17+   %13 = arith.addf %9, %12 : tensor<1024xf32>
18+   %14 = tt.splat %arg2 : (!tt.ptr<f32>) -> tensor<1024x!tt.ptr<f32>>
19+   %15 = tt.addptr %14, %4 : tensor<1024x!tt.ptr<f32>>, tensor<1024xi32>
20+   tt.store %15, %13, %6 {cache = 1 : i32, evict = 1 : i32} : tensor<1024x
21-   tt.return
22- }
```


Triton IR中的优化pass



先进编译实验室
Advanced Compiler

13

Rewrite-tensor-pointer Pass: 重写张量指针Pass

功能: 识别并重写那些涉及张量指针 (tensor pointers) 的操作, 以提高性能和内存使用效率。

执行命令: `triton-opt rewrite_before.ttir -triton-rewrite-tensor-pointer &> rewrite_after.ttir`

```
33  %20 = arith.extsi %arg6 : i32 to i64
34-  // CHECK-NOT: tt.make_tensor_ptr
35-  %21 = tt.make_tensor_ptr %arg0, [%18, %19], [%20, %c1_i64], [%16, %17]
36-  %22 = arith.muli %15, %c32_i32 : i32
37-  %23 = arith.extsi %arg4 : i32 to i64
38-  %24 = arith.extsi %arg7 : i32 to i64
39-  // CHECK-NOT: tt.make_tensor_ptr
40-  %25 = tt.make_tensor_ptr %arg1, [%19, %23], [%24, %c1_i64], [%17, %22]
41-  %26 = arith.addi %arg5, %c31_i32 : i32
42-  %27 = arith.divsi %26, %c32_i32 : i32
43-  %28 = arith.index_cast %27 : i32 to index
44-  %29:3 = scf.for %arg9 = %c0 to %28 step %c1 iter_args(%arg10 = %cst, %
45-  // CHECK: tt.load %{{.*}}, %{{.*}}, %{{.*}} {cache = 1 : i32, evict
```

```
33  %20 = arith.extsi %arg6 : i32 to i64
34+  %21 = arith.extsi %16 : i32 to i64
35+  %22 = arith.extsi %17 : i32 to i64
36+  %23 = arith.muli %15, %c32_i32 : i32
37+  %24 = arith.extsi %arg4 : i32 to i64
38+  %25 = arith.extsi %arg7 : i32 to i64
39+  %26 = arith.extsi %17 : i32 to i64
40+  %27 = arith.extsi %23 : i32 to i64
41+  %28 = arith.addi %arg5, %c31_i32 : i32
42+  %29 = arith.divsi %28, %c32_i32 : i32
43+  %30 = arith.index_cast %29 : i32 to index
44+  %31:5 = scf.for %arg9 = %c0 to %30 step %c1 iter_args(%arg10 = %cst, %ar
45+  %57 = tt.splat %arg0 : (!tt.ptr<f16>) -> tensor<128x32x!tt.ptr<f16>>
46+  %58 = tt.splat %arg11 : (i64) -> tensor<128xi64>
47+  %59 = tt.make_range {end = 128 : i32, start = 0 : i32} : tensor<128xi3
48+  %60 = arith.extsi %59 : tensor<128xi32> to tensor<128xi64>
49+  %61 = arith.addi %58, %60 : tensor<128xi64>
50+  %62 = tt.expand_dims %61 {axis = 1 : i32} : (tensor<128xi64>) -> tenso
51+  %63 = tt.splat %20 : (i64) -> tensor<128x1xi64>
52+  %64 = arith.muli %62, %63 : tensor<128x1xi64>
53+  %65 = tt.broadcast %64 : (tensor<128x1xi64>) -> tensor<128x32xi64>
54+  %66 = tt.addptr %57, %65 : tensor<128x32x!tt.ptr<f16>>, tensor<128x32x
55+  %67 = tt.splat %arg12 : (i64) -> tensor<32xi64>
56+  %68 = tt.make_range {end = 32 : i32, start = 0 : i32} : tensor<32xi32>
```



AdvancedCompiler

Tel: 13839830713